

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col
import matplotlib.pyplot as plt
import time
import os

spark = SparkSession.builder.appName("MeasureRunTime").getOrCreate()

for size in [100, 1000, 5000, 10000, 20000]:
    with open(f"data_{size}.csv", "w") as f:
        f.write("col\n")
        for i in range(size):
            f.write(f"{i % 10}\n") # Keeps grouping values small

for size in [100, 1000, 5000, 10000, 20000]:
    os.system(f"hdfs dfs -put -f data_{size}.csv /user/root/")

sizes = [100, 1000, 5000, 10000, 20000]
run_times = []

for size in sizes:
    df = spark.read.csv(f"/user/root/data_{size}.csv", header=True,
inferSchema=True)

    start_time = time.time()
    df.groupBy("col").avg("col").show()
    end_time = time.time()

    duration = end_time - start_time
    run_times.append(duration)

    print(f"File with {size} rows took {duration:.4f} seconds")
```

```
+---+-----+
|col|avg(col)|
+---+-----+
| 1|      1.0|
| 6|      6.0|
| 3|      3.0|
| 4|      4.0|
| 8|      8.0|
| 5|      5.0|
| 2|      2.0|
| 7|      7.0|
| 0|      0.0|
| 9|      9.0|
+---+-----+
```

File with 100 rows took 2.2430 seconds

col	avg(col)
1	1.0
6	6.0
3	3.0
4	4.0
8	8.0
5	5.0
2	2.0
7	7.0
0	0.0
9	9.0

File with 1000 rows took 0.6440 seconds

col	avg(col)
1	1.0
6	6.0
3	3.0
4	4.0
8	8.0
5	5.0
2	2.0
7	7.0
0	0.0
9	9.0

File with 5000 rows took 0.5281 seconds

col	avg(col)
1	1.0
6	6.0
3	3.0
4	4.0
8	8.0
5	5.0
2	2.0
7	7.0
0	0.0
9	9.0

File with 10000 rows took 0.4637 seconds

col	avg(col)
-----	----------

1	1.0
6	6.0
3	3.0
4	4.0
8	8.0
5	5.0
2	2.0
7	7.0
0	0.0
9	9.0

File with 20000 rows took 0.4885 seconds

```
plt.plot(sizes, run_times, marker='o')
plt.xlabel("File Size (Number of Rows)")
plt.ylabel("Run Time (Seconds)")
plt.title("Spark Runtime vs. File Size")
plt.grid(True)
plt.show()
```

